

SDS Platform

Description

SDS or "**Solid Distributed Systems**" is a platform for end-to-end infrastructure automation provided by Solid Distributed Systems SRL. Initial development started in 2023 and will be improved continuously.

The platform is composed from the shared infrastructure Kubernetes cluster and the controller instances/VMs which are hosted either in public cloud (AWS, Azure or GCP) or on bare metal with any virtualization platform.

From this core elements we deploy further Kubernetes clusters, cloud resources or ci/cd pipelines. Technologies used:

- Terragrunt and Terraform for cloud resources
- ArgoCD or FluxCD for Kubernetes configuration (current version supports only ArgoCD)
- Ansible for basic hosts configuration
- Tekton or GitHub Actions for Ci/CD (current version supports only Tekton inside Kubernetes)

Initial SDS platform bootstrap

Create the controller instances/VMs - step 1

Deploy 2 VMs or cloud instances with latest Ubuntu Server LTS version with the following requirements:

- access to internet to pull infrastructure code, docker images and OS patches (HTTPS - port 443)
- user that will be used on all hosts/VMs or instances.
- SUDO access for the user above with no password (see step 4.c bellow)
- Network access to all nodes and to entire infrastructure (ssh and https)
- 4vCPUs, 100 GB of fast storage and at least 8 GB of RAM
- AWS, Azure or GPC permissions to deploy remaining infrastructure.
- highly secured as this will contain all secrets, docker images and other critical platform data
- These hosts will be called "PLATFORM_NAME-controller" and "PLATFORM_NAME-controller-bkp" and will create/manage the entire platform infrastructure.
- These hosts will run ansible if the platform is hosted on bare-metal servers.
- These hosts will be configured with ansible in the same way, but only the first one will be active. The backup host will be switched manually in Disaster Recovery situations. These hosts will run core infrastructure services like NGINX (for services running on itself), Hashicorp Vault, Docker registry for air gaped environments, HAPROXY and Keepalived for on-prem Kubernetes clusters
- **IMPORTANT:** storage used by all services must be mounted as NFS volumes, so you must provide NFS volume for the controllers (NAS or in the cloud). This is required to perform the switch between controllers when needed.
- **IMPORTANT:** all commands and automations will run ONLY on the controller hosts.

Create the Kubernetes VMs for on-prem - step 2

If on-prem hosting is used you need 2 bare metal hosts with at least 18 CPU cores, 64 GB and 2TB storage each.

On these bare metal nodes you must setup a virtualization platform and the following VMs with latest Ubuntu Server LTS version as OS:

- 1 VMs on each bare metal host as Kubernetes control plane. hosts Each VM with 2vCPU, 50 GB of fast storage and 4GB RAM
- 2 VMs on each bare metal host as Kubernetes worker nodes. Each VM with 8vCPUs, 200 GB of storage and 16GB RAM
- Install on all 6 VMs latest Ubuntu Server LTS version and configure a platform wide used username (check step 4 bellow)

Configure DNS domain - step 3

- Configure public DNS domain that will contain all hosts and services for the platform.
- DNS service used must be updated programmatically from controller hosts and/or from external-dns service running inside Kubernetes cluster.
- Manual updates for DNS service are not supported, all DNS updates must be automatic.
- For current version only AWS Route53 is supported.

Initial controller configuration and SDS platform bootstrap - step 4

These manual steps bellow is to be refined and further automated with Ansible.

4.a - After installing latest Ubuntu Server LTS on both controller VMs, connect with PLATFORM_USERNAME and password to the first controller host and generate platform wide ssh key:

```
ssh-keygen -o -a 100 -t ed25519 -f ~/.ssh/id_rsa -C "USER_EMAIL_ADDRESS"
```

4.b - Add the public SSH key to the SDS github account and prepare parameter files, infrastructure code specifics (if any)

4.c - Configure sudo access with NOPASSW on all controller hosts and Kubernetes nodes: Scroll down till the end of the `/etc/sudoers` file and append the mentioned below line:

```
PLATFORM_USERNAME ALL=(ALL) NOPASSWD:ALL
```

4.d - Add platform wide SSH key to the second controller and to all Kubernetes cluster nodes:

```
ssh-copy-id PLATFORM_USERNAME@CONTROLLER_BKP_HOST
```

Also make sure you have to correct `~/.ssh/config` and `~/.ssh/id_rsa` on both controllers

4.e - On both controller hosts, in platform user home folder, clone the bootstrap repository:

```
git clone git@github.com:cbanciu667/sds-platform-bootstrap.git && cd sds-platform-bootstrap
```

4.f – Fill out the params file required for bootstrap according to examples (bash and ansible)

4.g - Run `./bootstrap.sh` to run ansible code and terraform with terragrunt. Some steps might be required to checked manually (for now let's encrypt certificate does not work).

Alternatively you may run the commands blocks from the script one after the other to monitor each action. This script will run ansible and trigger Kubespray or Terragrunt code for Kubernetes clusters creation.

4.g - Initialize and test the vault service which will provide secrets for Kubernetes clusters

- Run `docker-compose exec vault-server bash`
- Run `vault operator init` and save `vault init information` in a safe secret store
- Run `vault operator unseal UNSEAL_KEY1..3` (use 3 of the init tokens)
- Run `vault login ROOT_TOKEN`
- Run `vault audit enable file file_path=/vault/logs/audit.log`
- Put test secret `vault kv put secret/foo bar=precious`
- Get it with `vault kv get secret/foo`

4.h - Optionally, install VPN clients on both controllers to allow secure connectivity.

5 – Clone the SDS GitOps code (ArgoCD or FluxCD) and deploy Kubernetes resources

WORK IN PROGRESS:

- GitOps code
- Terragrunt/Terraform code

SOME REFERENCES:

<https://technekey.com/kubespray-advanced-configuration-for-a-production-cluster/>
https://schoolofdevops.github.io/ultimate-kubernetes-bootcamp/cluster_setup_kubespray/
<https://blog.devops.dev/multi-node-kubernetes-cluster-deployment-with-kubespray-and-ansible-c83c2c3c8f7f>
<https://github.com/gruntwork-io/terragrunt-infrastructure-live-example>