

High Available Kubernetes Cluster Setup using Kubespray

Kubespray is an *Ansible* based kubernetes provisioner. It helps us to setup a production grade, highly available and highly scalable Kubernetes cluster.

Prerequisites

Hardware Pre requisites

- 4 Nodes: Virtual/Physical Machines
- Memory: 2GB
- CPU: 1 Core
- Hard disk: 20GB available

Software Pre Requisites

On All Nodes

- Ubuntu 16.04 OS
- Python
- SSH Server
- Privileged user

On Ansible Control Node

- Ansible version 2.4 or greater

- Jinja

Networking Pre Requisites

- Internet access to download docker images and install softwares
- IPv4 Forwarding should be enabled
- Firewall should allow ssh access as well as ports required by Kubernetes. Internally open all the ports between node.

Architecture of a high available kubernetes cluster

Preparing the nodes

Run instructions in the section on `all nodes` in the cluster. This includes Ansible controller too.

Install Python

Ansible needs python to be installed on all the machines.

```
sudo apt update
sudo apt install python
```

Enable IPv4 Forwarding

On `all nodes`

Enable IPv4 forwarding by uncommenting the following line

```
echo "net.ipv4.ip_forward=1" >> /etc/sysctl.conf
```

Disable Swap

```
swapoff -a
```

Setup passwordless SSH between ansible controller and kubernetes nodes

On control node

Ansible uses passwordless ssh¹ to create the cluster. Let us see how to set it up from your *control node*.

Generate ssh keypair if not present already using the following command.

```
ssh-keygen -t rsa
```

Generating public/private rsa key pair.

Enter file in which to save the key (/home/ubuntu/.ssh/id_r

Enter passphrase (empty for no passphrase):

Enter same passphrase again:

Your identification has been saved in /home/ubuntu/.ssh/id_

Your public key has been saved in /home/ubuntu/.ssh/id_rsa.

The key fingerprint is:

```
SHA256:yC4Tl6RYc+saTPcLKFdGlTLOWOIuDgOlmy/NrMBnRxA ubuntu@n
```

The key's randomart image is:

```
+---[RSA 2048]-----+
|  E      ..          |
|  . o  +..          |
|  . +o*+o          |
|  . .o+Bo+          |
|  . .++ .X S          |
|+ +ooX .            |
| .=.OB.+ .          |
|  .o*= . .          |
|  .o. .            |
+-----[SHA256]-----+
```

Just leave the fields to defaults. This command will generate a public key and private key for you.

Copy over the public key to all nodes.

Example, assuming **ubuntu** as the user which has a privileged access on the node with ip address **10.10.1.101**,

```
ssh-copy-id ubuntu@10.10.1.101
```

This will copy our newly generated public key to the remote machine. After running this command you will be able to SSH into the machine directly without using a password. Replace *10.40.1.26* with your respective machine's IP.

e.g.

```
ssh ubuntu@10.10.1.101
```

Make sure to copy the public key to all kubernetes nodes. Replace username with the actual user on your system. If the above mentioned command fails, then copy your public key and paste it in the remote machine's `~/.ssh/authorized_keys` file.

e.g. (Only if ssh-copy-id fails)

```
cat ~/.ssh/id_rsa.pub  
ssh ubuntu@10.10.1.101  
vim ~/.ssh/authorized_keys  
# Paste the public key
```

Setup Ansible Control node and Kubespray

On control node

Set Locale

```
export LC_ALL="en_US.UTF-8"  
export LC_CTYPE="en_US.UTF-8"  
sudo dpkg-reconfigure locales
```

Do not select any other locale in the menu. Just press **(OK)** in the next two screens.

Setup kubespray

Kubespray is hosted on GitHub. Let us clone the [official repository](https://github.com/kubernetes-incubator/kubespray).

```
git clone https://github.com/kubernetes-incubator/kubespray
cd kubespray
```

Install Prerequisites

Install the python dependencies. **This step installs Ansible as well. You do not need to install Ansible separately.**

```
sudo apt install python-pip -y
sudo pip install -r requirements.txt
```

Set Remote User for Ansible

Add the following section in `ansible.cfg` file

```
remote_user=ubuntu
```

If the user you are going to connect is different, use that

instead.

Your *ansible.cfg* file should look like this.

```
[ssh_connection]
pipelining=True

ssh_args = -o ControlMaster=auto -o ControlPersist=30m -o C
#control_path = ~/.ssh/ansible-%%r@%%h:%%p

[defaults]
host_key_checking=False
gathering = smart
fact_caching = jsonfile
fact_caching_connection = /tmp
stdout_callback = skippy
library = ./library
callback_whitelist = profile_tasks
roles_path = roles:$VIRTUAL_ENV/usr/local/share/kubespray/r
deprecation_warnings=False
remote_user=ubuntu
```

Create Inventory

```
cp -rfp inventory/sample inventory/prod
```

where **prod** is the custom configuration name. Replace is with whatever name you would like to assign to the current cluster.

To build the inventory file, execute the inventory script

along with the IP addresses of our cluster as arguments

```
CONFIG_FILE=inventory/prod/hosts.ini python3 contrib/invent
```

Where replace the IP addresses (e.g. 10.10.1.101) with the actual IPs of your nodes

Once its run, you should see an inventory file generated which may look similar to below

```
file: inventory/prod/hosts.ini
```

```
[all]
```

```
node1      ansible_host=10.10.1.101 ip=10.10.1.101
```

```
node2      ansible_host=10.10.1.102 ip=10.10.1.102
```

```
node3      ansible_host=10.10.1.103 ip=10.10.1.103
```

```
node4      ansible_host=10.10.1.104 ip=10.10.1.104
```

```
[kube-master]
```

```
node1
```

```
node2
```

```
[kube-node]
```

```
node1
```

```
node2
```

```
node3
```

```
node4
```

```
[etcd]
```

```
node1
```

```
node2
```



```
node3
```

```
[k8s-cluster:children]
```

```
kube-node
```

```
kube-master
```

```
[calico-rr]
```

```
[vault]
```

```
node1
```

```
node2
```

```
node3
```

Customise Kubernetes Cluster Configs

There are two configs files in your inventroy directory's group_vars (e.g. inventory/prod/group_vars) viz.

- all.yml
- k8s-cluster.yml

Ansible is data driven, and most of the configurations of the cluster can be tweaked by changing the variable values from the above files.

Few of the configurations you may want to modify

```
file: inventory/prod/group_vars/k8s-cluster.yml
```

```
kubelet_max_pods: 100
```

```
cluster_name: prod
helm_enabled: true
```

Provisioning kubernetes cluster with kubespray

On control node

We are set to provision the cluster. Run the following `ansible-playbook` command to provision our Kubernetes cluster.

```
ansible-playbook -b -v -i inventory/prod/hosts.ini cluster.
```

Option `-i` = Inventory file path

Option `-b` = Become as root user

Option `-v` = Give verbose output

If you face this following error, while running `ansible-playbook` command, you can fix it by running following instructions

ERROR:

```
ERROR! Unexpected Exception, this is probably a bug: (crypt
```

FIX:

```
sudo pip install --upgrade pip
sudo pip uninstall cryptography
sudo pip install cryptography
ansible-playbook -b -v -i inventory/prod/hosts.ini cluster.
```

This Ansible run will take around 30 mins to complete.

Kubectl Configs

On kube master node

Once the cluster setup is done, copy the configuration and setup the permissions.

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Check the State of the Cluster

On the node where kubectl is setup

Let us check the state of the cluster by running,

```
kubectl cluster-info
```

```
Kubernetes master is running at https://10.10.1.101:6443
KubeDNS is running at https://10.10.1.101:6443/api/v1/names
```

To further debug and diagnose cluster problems, use 'kubect

```
kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
node1	Ready	master,node	21h	v1.9.0+coreos.0
node2	Ready	master,node	21h	v1.9.0+coreos.0
node3	Ready	node	21h	v1.9.0+coreos.0
node4	Ready	node	21h	v1.9.0+coreos.0

If you are able to see this, your cluster has been set up successfully.

¹ You can use private key / password instead of passwordless ssh. But it requires additional knowledge in using Ansible.

Access Kubernetes Cluster Remotely (Optional)

On your local machine

You could also install kubectl on your laptop/workstation. To learn how to install it for your OS, [refer to the procedure here](#).

e.g. To install **kubectl** on Ubuntu,

```
sudo apt-get update && sudo apt-get install -y apt-transport-https  
curl -s https://packages.cloud.google.com/apt/doc/apt-key.g
```

```
apt-key add -
```

```
sudo touch /etc/apt/sources.list.d/kubernetes.list
```

```
echo "deb http://apt.kubernetes.io/ kubernetes-xenial main"
```

```
sudo apt-get update
```

```
sudo apt-get install -y kubectl
```

Copy kubernetes config to your local machine

Copy kubeconfig file to your local machine

```
mkdir ~/.kube
```

```
scp -r ubuntu@MASTER_HOST_IP:/etc/kubernetes/admin.conf ~/
```

```
kubectl get nodes
```

Deploy Kubernetes Objects

Since its a new cluster, which is differnt than what you have created with kubeadm earlier, or if this is the first time you are creating a kubernetes cluster with kubespray as part of **Advanced Workshop**, you need to deploy services which have been covered as part of the previous

topics.

In order to do that, use the following commands on the node where you have configured kubectl

```
git clone https://github.com/schoolofdevops/k8s-code.git
```

```
cd k8s-code/projects/instavote
```

```
kubectl apply -f instavote-ns.yaml
```

```
kubectl apply -f prod/
```

Switch to **instavote** namespace and validate,

```
kubectl config set-context $(kubectl config current-context
```

```
kubectl get pods,deploy,svc
```

where,

- `--cluster=prod` : prod is the cluster name you created earlier. If not, use the correct name of the cluster (`kubectl config view`)
- `--user=admin-prod` : is the admin user created by default while installing with kubespray
- `--namespace=instavote` : the namespace you just created to deploy instavote app stack

[sample output]

```
$ kubectl get pods,deploy,svc
```

NAME	READY	STATUS	RESTARTS
pod/db-66496667c9-qggzd	1/1	Running	0
pod/redis-6555998885-4k5cr	1/1	Running	0
pod/redis-6555998885-fb8rk	1/1	Running	0
pod/result-5c7569bcb7-4fptr	1/1	Running	0
pod/result-5c7569bcb7-s4rdx	1/1	Running	0
pod/vote-5d88d47fc8-gbzbq	1/1	Running	0
pod/vote-5d88d47fc8-q4vj6	1/1	Running	0
pod/worker-7c98c96fb4-7tzzw	1/1	Running	0

NAME	DESIRED	CURRENT	UP-TO-DA
deployment.extensions/db	1	1	1
deployment.extensions/redis	2	2	2
deployment.extensions/result	2	2	2
deployment.extensions/vote	2	2	2
deployment.extensions/worker	1	1	1

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP
service/db	ClusterIP	10.233.16.207	<none>
service/redis	ClusterIP	10.233.14.61	<none>
service/result	NodePort	10.233.22.10	<none>
service/vote	NodePort	10.233.19.111	<none>

References

- [Installing Kubernetes On Premises/On Cloud with Kubespray](#)
- [Kubespray on Github](#)

