

Kubespray Advanced Configuration For A Production Cluster

/ [Kubernetes](#), [tech](#)

In this post, we will see some advanced options provided by kubespray for production-grade cluster creation. In most places on the internet, you will see the basic instructions to create a simple(Hello World) Kubernetes cluster using a bunch of virtual machines. However, kubespray offer much more than that. For example:

1. You can install several add-ons required in the daily operations of the cluster.
2. Kubespray provides capabilities to expose control plan nodes via pre-existing load balancers.

We will explore some of the critical configurations. However, there would always be more in the official documents.

You might want to see [this page](#) if you intend to create a simple cluster that works fine for development or learning.

If you are unfamiliar with load-balancing, consider reading the basics [here](#) to make more sense of the information provided in this post.

The entire procedure described here is automated and available; check [this page](#) for details.

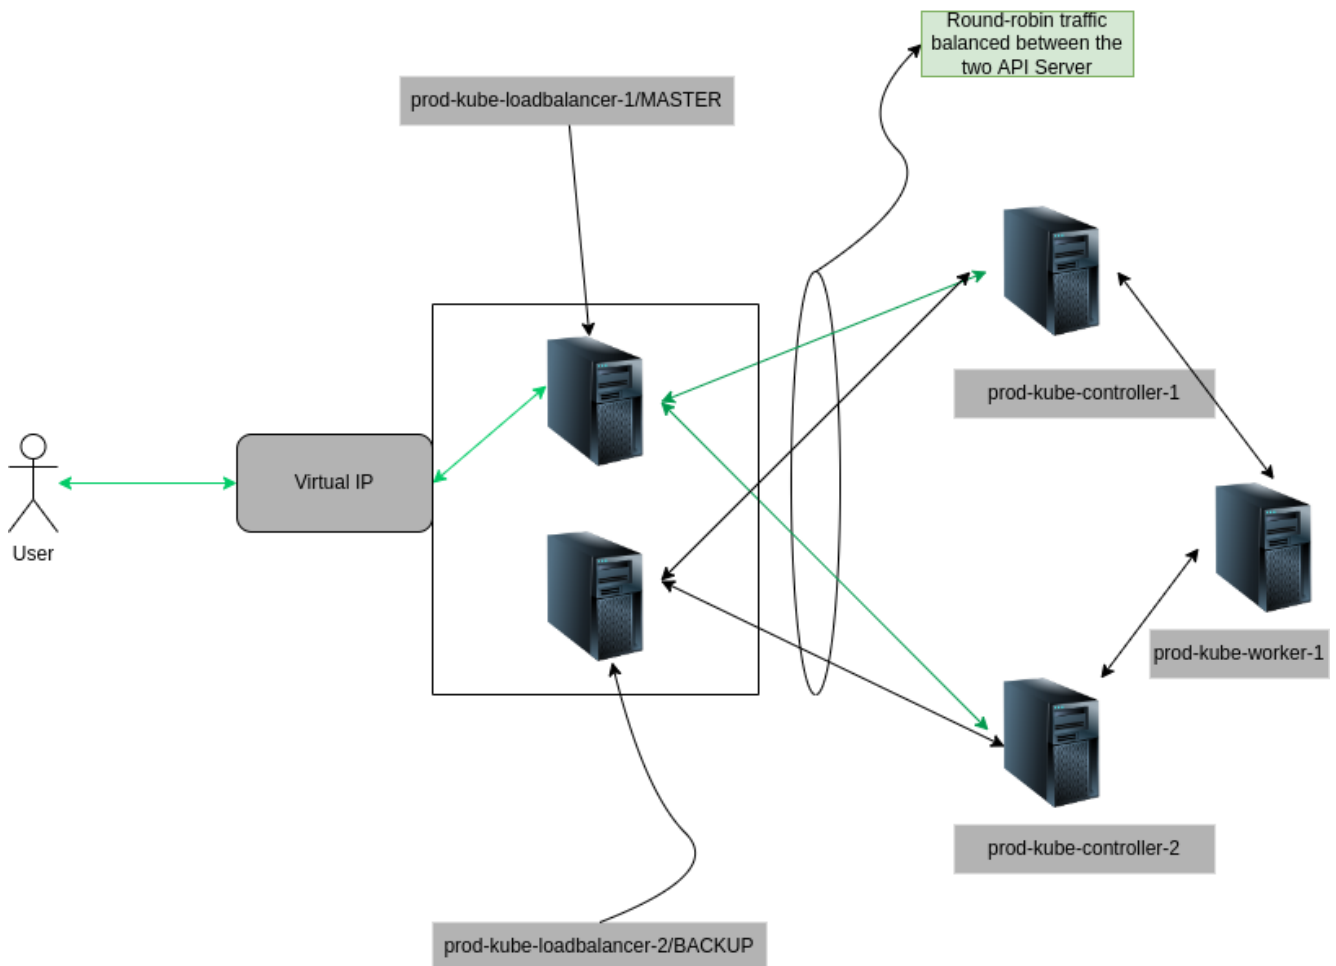
Use case:

The following is the objective we will achieve in this post. These actions will make the cluster resilient to a single point of failure and provide several add-ons ready to use after cluster instantiation.

- 1. Create a High-available cluster, using multiple(two) Kube-API servers and one worker node(You can have more).
- 2. Create multiple frontend external load balancers nodes(2).
- 3. Configure the external load balancers, and expose the Kube-API server via Virtual IP.
- 4. Using kubespray addons, Install helm, Kubernetes dashboard, metric server, and a few more addons.

The Setup: I have already created the following virtual machines

Name	MAC address	Protocol	Address
vnet20	52:54:00:97:cb:dd	ipv4	192.168.122.1
vnet21	52:54:00:19:d9:6e	ipv4	192.168.122.1
vnet22	52:54:00:90:40:35	ipv4	192.168.122.2
vnet23	52:54:00:ca:4c:d1	ipv4	192.168.122.2



Traffic path in Green, when LB-1 is VRRP Master

Start of the Procedure:

Step-1: Install git and Pip3

```
sudo apt update && sudo apt-get -y install python3-pip git
```

Step-2: Clone The Git Repo

```
git clone https://github.com/kubernetes-sigs/kubespray.git
```

Step-3: Install the requirements

```
cd kubespray/  
sudo pip3 install -r requirements.txt
```

Step-4: Copy the blank inventory template

```
cp -rfp inventory/sample inventory/mycluster
```

step-5: Make a list/array of hostname and IP for all controller and worker nodes

```
declare -a IPS=(controller-vm-hostname-1,IP-1 controller-vm
```

```
declare -a IPS=(prod-kube-controller-1,192.168.122.105 prod
```

Step-6: Build the inventory

```
CONFIG_FILE=inventory/mycluster/hosts.yaml python3 contrib/
```

```
CONFIG_FILE=inventory/mycluster/hosts.yaml python3 contrib/
```

```
DEBUG: Adding group all
```

```
DEBUG: Adding group kube_control_plane
```

```
DEBUG: Adding group kube_node
DEBUG: Adding group etcd
DEBUG: Adding group k8s_cluster
DEBUG: Adding group calico_rr
DEBUG: adding host rod-kube-controller-1 to group all
DEBUG: adding host prod-kube-controller-2 to group all
DEBUG: adding host prod-kube-worker-1 to group all
DEBUG: adding host rod-kube-controller-1 to group etcd
DEBUG: adding host prod-kube-controller-2 to group etcd
DEBUG: adding host prod-kube-worker-1 to group etcd
DEBUG: adding host rod-kube-controller-1 to group kube_cont
DEBUG: adding host prod-kube-controller-2 to group kube_con
DEBUG: adding host rod-kube-controller-1 to group kube_node
DEBUG: adding host prod-kube-controller-2 to group kube_nod
DEBUG: adding host prod-kube-worker-1 to group kube_node
```

Step-7: Understand, Validate and visualize the generated inventory file

In the inventory file, there are three host groups created.

1. **kube_control_plane:** the nodes under this host group will be configured as controller nodes(API server will run here)
2. **kube_node:** The nodes under this host group will be configured as worker nodes. (payload nodes)
3. **etcd:** The nodes under this host group will be running the etcd database for the cluster. The number of nodes under this host group must be odd. To meet, within the etcd quorum.

```
cat inventory/mycluster/hosts.yaml
```

all:

hosts:

rod-kube-controller-1:

ansible_host: 192.168.122.105

ip: 192.168.122.105

access_ip: 192.168.122.105

prod-kube-controller-2:

ansible_host: 192.168.122.137

ip: 192.168.122.137

access_ip: 192.168.122.137

prod-kube-worker-1:

ansible_host: 192.168.122.253

ip: 192.168.122.253

access_ip: 192.168.122.253

children:

kube_control_plane:

hosts:

rod-kube-controller-1:

prod-kube-controller-2:

kube_node:

hosts:

rod-kube-controller-1:

prod-kube-controller-2:

prod-kube-worker-1:

etcd:

hosts:

rod-kube-controller-1:

prod-kube-controller-2:

prod-kube-worker-1:

k8s_cluster:

children:

kube_control_plane:

kube_node:

calico_rr:

```
hosts: {}
```

Now the inventory file is reNext, we next; we can work towards more advanced configuration options. Most of these options are mutually exclusive from each other, so if you skip anyone, likely, it won't trouble you.

File-1: Tweaking the cluster configuration

The first file we would want to tweak is **inventory/mycluster/group_vars/k8s_cluster/k8s-cluster.yml**. These are the bare minimum configurations I would like to highlight. Feel free to explore the rest of the file and tweak it if required.

The last variable, **"supplementary_addresses_in_ssl_keys,"** is critical for load balancers to work; in this step, we tell kubeadm to add VIP while creating the certificates.

```
kube_version: v1.24.3
```

```
container_manager: containerd
```

`cluster_name: prod.local`

`kube_network_plugin: calico`

`k8s_image_pull_policy: IfNotPresent`

`kubernetes_audit: true`

`kube_encrypt_secret_data: true`

`kubeconfig_localhost: true`

`kubectl_localhost: true`

`supplementary_addresses_in_ssl_keys: [10.0.0.1, 10.0.0.2, 1`

File-2: adding the external load balancer VIP

In this section, we will provide the external load balancer Virtual IP and port in

inventory/mycluster/group_vars/all/all.yml

```
loadbalancer_apiserver:  
  address: "192.168.122.211"  
  port: "8443"
```

File-3: Manage the addons

I am limiting this section to the dashboard, registry, helm, metrics server, an ingress controller. However, many more add-ons are present in the `inventory/mycluster/group_vars/k8s_cluster/addons.yml` file—for example, metallb, cert-manager, etc. Consider going through the file and choosing what you need.

```
dashboard_enabled: true
```

```
helm_enabled: true
```

```
registry_enabled: true  
registry_namespace: kube-system  
registry_storage_class: ""  
registry_disk_size: "10Gi"
```

metrics_server_enabled: true
metrics_server_container_port: 4443
metrics_server_kubelet_insecure_tls: true
metrics_server_metric_resolution: 15s
metrics_server_kubelet_preferred_address_types: "InternalIP"

ingress_nginx_enabled: true
ingress_nginx_host_network: true
ingress_publish_status_address: ""
ingress_nginx_nodeselector:
 kubernetes.io/os: "linux"
ingress_nginx_tolerations:
 - key: "node-role.kubernetes.io/master"
 operator: "Equal"
 value: ""
 effect: "NoSchedule"
 - key: "node-role.kubernetes.io/control-plane"
 operator: "Equal"
 value: ""
 effect: "NoSchedule"
ingress_nginx_namespace: "ingress-nginx"
ingress_nginx_insecure_port: 80
ingress_nginx_secure_port: 443
ingress_nginx_class: nginx

Now we have desired configuration done, and all add-ons are set up. We will now set up the load balancers. The kubespray does not cover this process. We will use keepalived and haproxy for the configuration. In this post, we will superficially go through the installation. However, You can read more about load balancers in great detail [HERE](#). If you already have load balancers configured, you may skip all the following steps(LB-step-xx)

LB-step-1: Install keepalived and haproxy by running the following commands on all load balancers

```
sudo apt-get update
sudo apt-get upgrade
sudo apt install haproxy -y
```

```
sudo apt install keepalived -y
```

LB-Step-2: Create the keepalived config file on all load balancer nodes with minor changes described below

Add the following text in

/etc/keepalived/keepalived.conf present on all load balancer nodes. Note that I have added the Virtual IP here. Copy the same content to all the load balancer nodes, **except** change the state to BACKUP and decrease the priority to 254.

IMPORTANT NOTE: My load balancer nodes are running

ubuntu22.04, which has 'enp1s0' as a default interface. If you have some other interface, you might need to replace it with the right one.

```
global_defs {
router_id LVS_DEVEL
script_user root
enable_script_security
}

vrrp_script check_apiserver {
    script "/etc/keepalived/check_apiserver.sh"
    interval 3
    weight -2
    fall 10
    rise 2
}

vrrp_instance VI_1 {
    state MASTER
    interface enp1s0
    virtual_router_id 51
    priority 255
    authentication {
        auth_type PASS
        auth_pass mypass
    }
    virtual_ipaddress {
        192.168.122.211/24
    }
    track_script {
        check_apiserver
    }
    notify_master "/etc/keepalived/status_capture.sh MASTER
```

```
    notify_backup "/etc/keepalived/status_capture.sh BACKUP"
    notify_fault  "/etc/keepalived/status_capture.sh FAULT"
}
```

LB-Step-3: Create an identical Keepalived check script on all the load balancer node

Add the following content to
/etc/keepalived/check_apiserver.sh

```
errorExit() {
    echo "*** $*" 1>&2
    exit 1
}

curl --silent --max-time 2 --insecure https://localhost:844

if ip addr | grep -q 192.168.122.211; then
    curl --silent --max-time 2 --insecure https://192.168.1
fi
```

LB-Step-4: Create an identical keepalived notify script on both load balancer node

Create a notify script at
/etc/keepalived/status_capture.sh; this is optional but

very helpful in real-world scenarios to trigger actions based on state change.

```
#  
echo "$(date): The loadbalancer instance running on $(hostname)  
chmod 755 /tmp/load-balancer-status || true
```

LB-Step-5: Create an identical haproxy config file on both load balancer node

Create the haproxy config file at
/etc/haproxy/haproxy.cfg

```
defaults  
    mode tcp  
    timeout connect 10s  
    timeout client 30s  
    timeout server 30s  
  
frontend apiserver  
    bind *:8443  
    mode tcp  
    option tcplog  
    log 127.0.0.1 local0  
    default_backend apiserver  
  
backend apiserver  
    option httpchk GET /healthz  
    http-check expect status 200  
    mode tcp  
    option ssl-hello-chk
```

```
balance        roundrobin
server kube-controller-1 192.168.122.105:6443 check
server kube-controller-2 192.168.122.137:6443 check
```

LB-Step-6: Make the keepalived check and notify script executable on both load balancer nodes

```
sudo chmod u+x /etc/keepalived/check_apiserver.sh
sudo chmod u+x /etc/keepalived/status_capture.sh
```

LB-Step-7: Enable non-local binding on both load balancer nodes

```
echo 'net.ipv4.ip_nonlocal_bind=1'|sudo tee -a /etc/sysctl.
sudo sysctl -p
```

LB-Step-8: Enable keepalived and haproxy on both load balancer nodes

```
sudo service keepalived start
sudo service keepalived status
sudo service haproxy start
sudo service haproxy status
```

We have created the load balancer nodes, and the kubespray configuration is ready. Trigger the playbook for

cluster creation.

Trigger the cluster creation

```
ansible-playbook -i inventory/mycluster/hosts.yaml --become
```

Once the playbook completes, you will notice zero failure, indicating that whatever we have done so far is working.

```
PLAY RECAP *****
localhost                        : ok=3      changed=0    unreachable
prod-kube-controller-1          : ok=813   changed=136  unreachable
prod-kube-controller-2          : ok=711   changed=125  unreachable
prod-kube-worker-1              : ok=560   changed=92   unreachable
```

Locate the kubeconfig file

Remember, we set up a variable called **kubeconfig_localhost** in earlier steps. That caused the kubeconfig to be copied to **inventory/mycluster/artifacts/admin.conf**

```
ls -lrt inventory/mycluster/artifacts/admin.conf
-rw----- 1 technekey technekey 5683 Jul 21 14:06 inventor
```

Checking the addons and cluster state


```
kubectl get pod -A --kubeconfig inventory/mycluster/artifac
```

NAMESPACE	NAME
ingress-nginx	ingress-nginx-controller-59g57
ingress-nginx	ingress-nginx-controller-bnxvf
ingress-nginx	ingress-nginx-controller-djv4p
kube-system	calico-node-5pxff
kube-system	calico-node-gnkhj
kube-system	calico-node-k6xjs
kube-system	coredns-74d6c5659f-d9n87
kube-system	coredns-74d6c5659f-pqd4n
kube-system	dns-autoscaler-59b8867c86-z7dhx
kube-system	kube-apiserver-prod-kube-controller-1
kube-system	kube-apiserver-prod-kube-controller-2
kube-system	kube-controller-manager-prod-kube-controlle
kube-system	kube-controller-manager-prod-kube-controlle
kube-system	kube-proxy-6h96q
kube-system	kube-proxy-9hpg2
kube-system	kube-proxy-q72x6
kube-system	kube-scheduler-prod-kube-controller-1
kube-system	kube-scheduler-prod-kube-controller-2
kube-system	kubernetes-dashboard-55bf5db569-m4sqm
kube-system	kubernetes-metrics-scraper-84bbbc8b75-28zr9
kube-system	metrics-server-68b8967c9f-4fprw
kube-system	nodelocaldns-9mlxt
kube-system	nodelocaldns-k2qqg
kube-system	nodelocaldns-zrbg6
kube-system	registry-mw4b6

```
kubectl get svc -n kube-system --kubeconfig inventory/myclu
```

NAME	TYPE	CLUSTER-IP	EXT
coredns	ClusterIP	10.233.0.3	<no

dashboard-metrics-scraper	ClusterIP	10.233.58.178	<no
kubernetes-dashboard	ClusterIP	10.233.57.157	<no
metrics-server	ClusterIP	10.233.24.49	<no
registry	ClusterIP	10.233.18.18	<no

kubectl top node --kubeconfig inventory/mycluster/artifacts

NAME	CPU(cores)	CPU%	MEMORY(bytes)
prod-kube-controller-1	184m	10%	1503Mi
prod-kube-controller-2	176m	9%	1407Mi
prod-kube-worker-1	114m	6%	1197Mi

kubectl get pod -n ingress-nginx --kubeconfig inventory/my

NAME	READY	STATUS	RESTARTS
ingress-nginx-controller-59g57	1/1	Running	0
ingress-nginx-controller-bnxvf	1/1	Running	0
ingress-nginx-controller-djv4p	1/1	Running	0

Related reading and references:

- <https://github.com/kubernetes/kubeadm/blob/main/docs/ha-considerations.md>
- <https://github.com/kubernetes-sigs/kubespray/blob/master/docs/ha-mode.md>
- <https://technekey.com/ha-kubernetes-cluster-using-keepalived-and-haproxy/>